

IBM Watson Seminar

String Kernels and Cluster Kernels for Protein Classification


Christina Leslie

Department of Computer Science

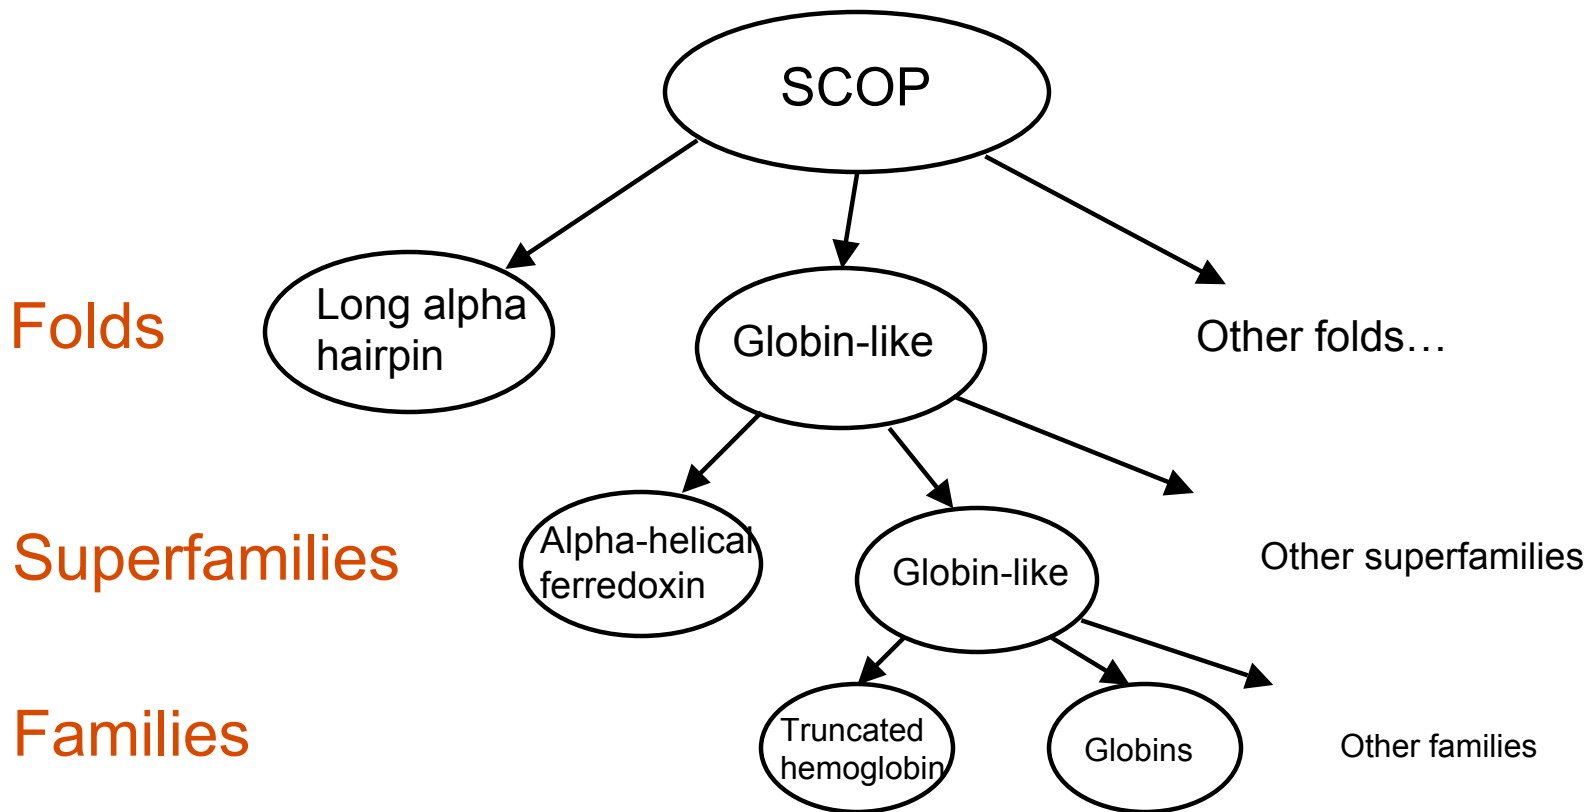
Columbia University

Protein Sequence Classification

- Protein represented by sequence of amino acids, encoded by a gene
- Easy to sequence proteins, difficult to obtain structure
- *Classification Problem*: Learn how to classify protein sequence data into families and superfamilies defined by structure/function relationships



Structural Hierarchy of Proteins



- *Remote homologs*: sequences that belong to the same superfamily but not the same family – remote evolutionary relationship
- Structure and function conserved, though sequence similarity can be low

Learning Problem

- Use *discriminative supervised learning* approach to classify protein sequences into structurally-related families, superfamilies
- Labeled training data: proteins with *known structure*, positive (+) if example belongs to a family or superfamily, negative (-) otherwise
- Focus on *remote homology detection*

Approach: Support vector machines (SVMs) with new *string kernels* based on inexact string matching

Labeled Training
Sequences



Classification
Rule

Learning Algorithm

Beyond Classification: Protein Ranking

- *Ranking problem*: given query protein sequence, return ranked list of similar proteins from sequence database
- Limitations of classification framework
 - Small amount of labeled data (proteins with known structure), huge unlabeled databases
 - Missing classes: undiscovered structures
- Good local similarity scores, based on heuristic alignment: BLAST, PSI-BLAST

Approach: Use new *semi-supervised learning* methods – training on labeled and unlabeled data – to improve ranking performance

Outline

1. Protein classification: Mismatch kernel
 - SVMs and kernel methods
 - Inexact matching through mismatches
 - Efficient kernel computation, fast prediction
2. Experimental results on SCOP dataset
3. Other models for inexact matching
 - Kernels from gaps, substitutions, wildcards
4. Cluster kernels: Semi-supervised methods
 - Using unlabeled data to change the kernel

Support Vector Machine (SVM) Classifiers

- Training examples mapped to (usually high-dimensional) feature space by a *feature map*


$$F(x) = (F_1(x), \dots, F_N(x))$$

- Learn linear classifier in feature space

$$f(x) = \mathbf{w}^T F(x) + b$$

by solving optimization problem:
trade-off between maximizing *geometric margin* and minimizing margin violations

- Large margin gives good generalization performance, even in high dimensions



Kernels for Discrete Objects

- **Kernel trick:** To train an SVM, can use *kernel* rather than explicit feature map
- Can define kernels for sequences, graphs, other *discrete* objects:

$$\{ \text{sequences} \} \xrightarrow{F} \mathbb{R}^N$$

For sequences x, y , feature map F , kernel value is inner product in feature space

$$K(x, y) = \langle F(x), F(y) \rangle$$

- Original string kernels [Watkins, Haussler, later Lodhi *et al.*] require quadratic time in sequence length, $O(|x| |y|)$, to compute each kernel value $K(x, y)$


String Kernels for Biosequences

- We'll define new fast *string kernels* for biological sequence data
 - Biologically-inspired underlying feature map
 - Kernels scale linearly with sequence length, $O(c_K(|x| + |y|))$ to compute
 - Protein classification performance competitive with best available methods
 - Mismatches for *inexact sequence matching* (other models later)

Spectrum-based Feature Map

- Idea: feature map based on *spectrum* of a sequence
 - The k-spectrum of a sequence is the set of all k-length contiguous subsequences that it contains
 - Feature map is indexed by all possible k-length subsequences (“k-mers”) from the alphabet of amino acids
 - Dimension of feature space = $|\Sigma|^k$ ($|\Sigma| = 20$ for amino acids)

AKQDYYYEI



AKQ

KQD

QDY

DYY

YYY

YYY


YYE

YEI

k-Spectrum Feature Map

- Feature map for k-spectrum with no mismatches:


For sequence x , $F_{(k)}(x) = (F_t(x))_{\{k\text{-mers } t\}}$,
where $F_t(x) = \text{\#occurrences of } t \text{ in } x$



C. Leslie, E. Eskin, and W. Noble, *The Spectrum Kernel: A String Kernel for SVM Protein Classification*. Pacific Symposium on Biocomputing, 2002.

Inexact Matching through Mismatches

- For k-mer s , the *mismatch neighborhood* $N_{(k,m)}(s)$ is the set of all k-mers t within m mismatches from s
- Size of mismatch neighborhood is $O(|\Sigma|^m k^m)$



(k,m)-Mismatch Feature Map

- Feature map for k-spectrum, allowing m mismatches:

For a k-mer s , $F_{(k,m)}(s) = (F_t(s))_{\{k\text{-mers } t\}}$,
where $F_t(s) = 1$ if t is in neighborhood $N_{(k,m)}(s)$,
 $F_t(s) = 0$ otherwise


AKQ \longrightarrow (0 , ... , 1 , ... , 1 , ... , 1 , ... , 1 , ... , 0)
 AAQ AKQ DKQ EKQ

- Extend additively to longer sequences x by summing over all k-mers s in x

C. Leslie, E. Eskin, J. Weston and W. Noble, *Mismatch String Kernels for SVM Protein Classification*. Neural Information Processing Systems 2002.

Computing the (k,m)-Mismatch Kernel

- Use *mismatch tree* to organize lexical traversal of all instances of k-mers (with mismatches) in the training data
 - Each path down to a leaf corresponds to a coordinate in feature map
 - Kernel values for all training sequences updated at each leaf node
 - Depth-first traversal can be accomplished with recursive function




Computing the Kernel for Pair of Sequences

- Traversal of trie for $k=3$, $m=1$

x : ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓
 E A D L A L G K A V F

y : A D L A L G A D Q V F N G
 ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑



Computing the Kernel for Pair of Sequences


- Traversal of trie for $k=3$, $m=1$

x : **E**AD**L**AL**G**KAV**F**

 ↓ ↓ ↓

y : **A**D**L**AL**G**AD**Q**V**F**NG

 ↑ ↑ ↑




Computing the Kernel for Pair of Sequences

- Traversal of trie for $k=3$, $m=1$

x : EADLALGKAVF
 ↓
 y : ADLALGADQVFNG
 ↑ ↑

Update kernel value for $K(x, y)$ by adding contribution for feature ADL




Scales linearly with length,
 $O(k^{m+1}|\Sigma|^m(|x|+|y|))$

SVM Solution

- Linear classifier defined in feature space by
$$f(x) = \mathbf{w}^T F(x) + b$$
where $\text{sign}(f(x))$ gives prediction
- SVM solution gives normal vector

$$\mathbf{w} = \sum_i y_i \sum_i F(x_i)$$

as a linear combination of *support vectors*, involving weights \sum_i and labels y_i



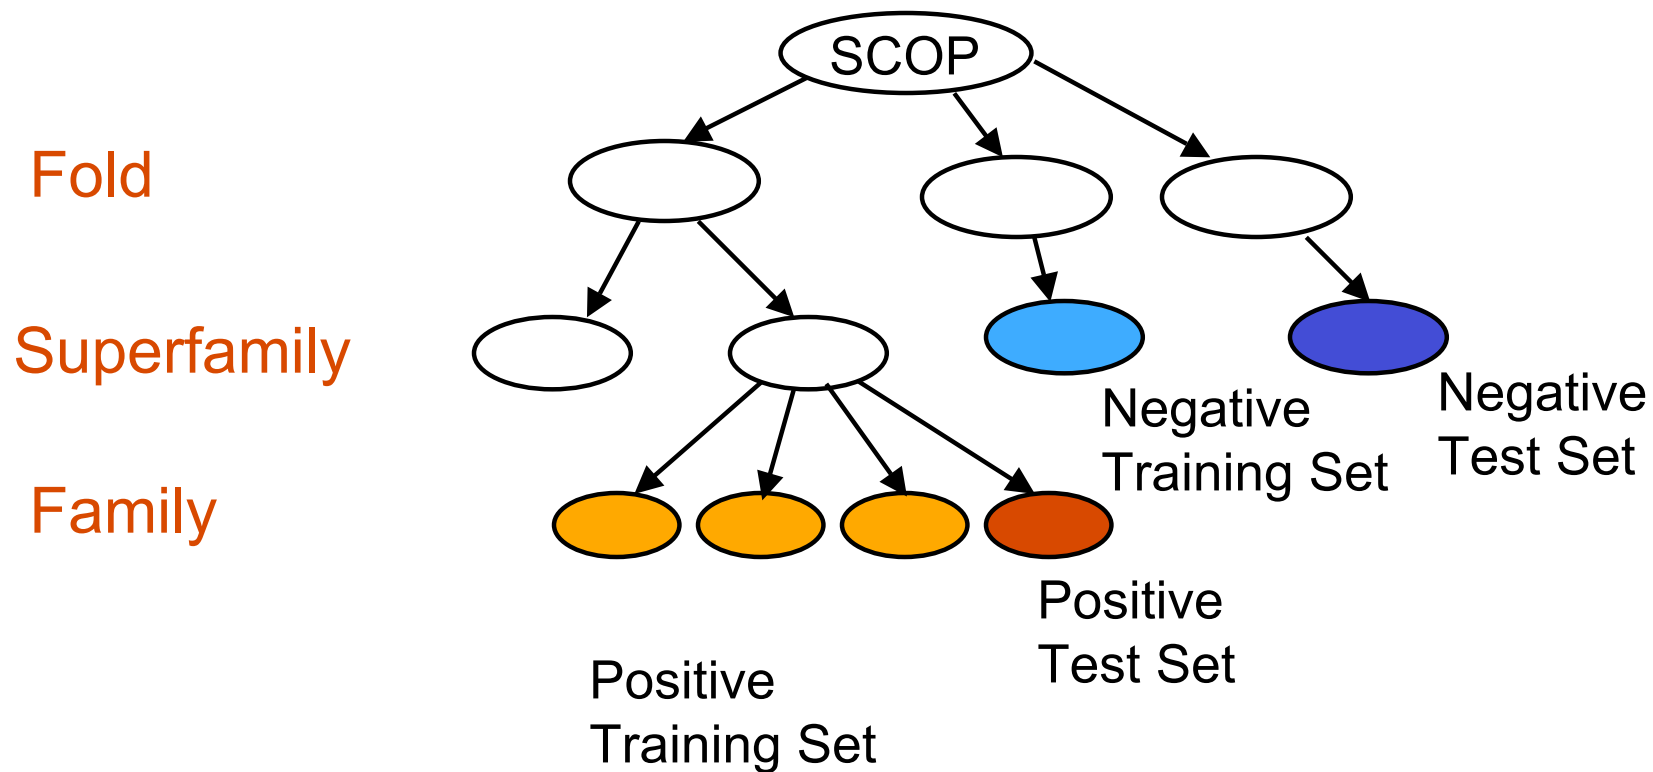
Fast prediction

- SVM training determines subset of training sequences corresponding to *support vector sequences* and their weights: (x_i, α_i)
- Linear decision rule in feature space:
$$f(x) = \sum_i y_i \alpha_i \sum F(x_i), F(x) + b$$
- $F(x)$ is sum of feature vectors $F(s)$ for k-mers s in x
 - Precompute per k-mer scores for classifier
 - Test sequences can be classified in *linear time* via lookup of k-mers

Outline

1. Protein classification: Mismatch kernel
 - SVMs and kernel methods
 - Inexact matching through mismatches
 - Efficient kernel computation, fast prediction
2. Experimental results on SCOP dataset
3. Other models for inexact matching
 - Kernels from gaps, substitutions, wildcards
4. Cluster kernels: Semi-supervised methods
 - Using unlabeled data to change the kernel

SCOP Experiments




- Tested with experiments on SCOP dataset from Jaakkola *et al.*
- Experiments designed to ask: Could the method discover a new family of a known superfamily?

SCOP Experiments

- 160 experiments for 33 target families from 16 superfamilies
- Compared results against
 - SVM-Fisher (HMM-based kernel)
 - SAM-T98 (profile HMM)
 - PSI-BLAST (heuristic alignment-based method)
- *ROC scores*: area under the graph of true positives as a function of false positives, scaled so that both axes vary between 0 and 1

Results Across All Target Families



Background on Fisher-SVM

- Previous solution [Jaakkola, Diekhans, Haussler]:
 - Use positive examples to train profile HMM, (M_+, θ_0)
 - For each training example x , *Fisher score* is gradient of log-likelihood score for x given M_+ (evaluated at θ_0)
$$x \longrightarrow \frac{\partial}{\partial \theta} \log P(x \mid M_+, \theta)$$
- Method relies on generative model
 - Requires large amount of data or sophisticated priors to train M_+
 - Expensive: dynamic programming (quadratic in sequence length) – for each sequence x , forward-backward algorithm to compute features

Aside: Connection with Fisher Kernel

- Consider order $k-1$ Markov chain model for positive sequences, with parameters

$$\varphi^t_{s_1 \dots s_{k-1}} = P(x_j = t \mid x_{j-k+1} \dots x_{j-1} = s_1 \dots s_{k-1})$$

- Corresponding Fisher coordinate for x is

$$\begin{aligned} & (\# \text{occurrences of } s_1 \dots s_{k-1} t \text{ in } x) / \varphi^t_{s_1 \dots s_{k-1}} \\ & \quad \varphi(s_1 \dots s_{k-1}) \end{aligned}$$

- Fisher kernel for Markov chain model similar to k -spectrum kernel

Interpretation of Mismatch-SVM Classifier


- Rank features by $|w_i|$, associate to +/- class by $\text{sign}(w_i)$
- Top positively-weighted k-mer features learned by SVM map to *conserved regions* in the *multiple alignment* of positive training sequences

```
>saq.---ksp..aelksifek..yaa keg....dpnqlsk...eel.....  
..kqliqaef.....p.....  
>snklh-----..----fafrl..yd-ldk....d-ekisr...del.....  
..lqvlrmmv.....gvnisdeqlgsi  
>maa..pldqai..gllvatfhk..ysgkeg....dknslsk...gel.....  
..keliqkelti.....g.....  
>sfggf-----..----kvfd-..---edg....d-gyisa...rel.....  
..qmvlgk-1.....g.....  
>snklh-----..----fafrl..yd-ldk....d-dkisr...del.....  
..lqvlrmmv.....gvnisdeqlgsi
```

```
>sll.....k.....g.....prtlddl.  
...fqeldkn...gdgevSFEEFQvlvkkisq.....  
>---.....-.....-.....---adrt.  
...iqeadqd...gdsaiSFTEFVkv1----ekvdv.  
>pk1.....k.....d.....ae-iagl.  
...medldrn...kdqeVNFQEyvtflgalamiynea.  
>---.....-fsegs.e.....idrvekm.  
...ivsvdsn...rdGRVDFFEFkdmm----rsvlv.  
>---.....-.....-.....---adrt.  
...iqeadqd...gdsaiSFTEFVkv1----ekvdv.
```


Interpretation of Mismatch-SVM Classifier

- Rank features by $|w_i|$, associate to +/- class by $\text{sign}(w_i)$
- Top positively-weighted k-mer features learned by SVM map to *conserved regions* in the *multiple alignment* of positive training sequences



Interpretation of Mismatch-SVM Classifier

- Rank features by $|w_i|$, associate to +/- class by sign
- Top positively-weighted k-mer features learned by SVM map to *conserved regions* in the *multiple alignment* of positive training sequences



Advantages of Mismatch-SVM

- Mismatch-SVM performs as well as SVM-Fisher but avoids computational expense, training difficulties of profile HMM
- Advantages of string kernel:
 - *Efficient computation*: scales linearly with sequence length
 - *Fast prediction*: classify test sequences in linear time
 - *Interpretation* of learned classifier
 - *General approach* for biosequence data, does not rely on alignment or generative model

Outline

1. Protein classification: Mismatch kernel
 - SVMs and kernel methods
 - Inexact matching through mismatches
 - Efficient kernel computation, fast prediction
2. Experimental results on SCOP dataset
3. Other models for inexact matching
 - Kernels from gaps, substitutions, wildcards
4. Cluster kernels: Semi-supervised methods
 - Using unlabeled data to change the kernel


Other Fast(er) Kernels for Inexact Matching

- Mismatch kernel is linear in sequence length, but constant $c_K = k^{m+1}|\Sigma|^m$ depends on alphabet size
- Other models for inexact matching can achieve $O(c_K(|x| + |y|))$ with c_K independent of $|\Sigma|$
 - Restricted gaps
 - Probabilistic substitutions
 - Wildcards

C. Leslie and R. Kuang, *Fast String Kernels for Inexact String Matching*. To appear, COLT/KW 2003.

Inexact Matching through Gaps

- For g-mer s , the *gapped match set* $G_{(g,k)}(s)$ consists of all k-mers t that occur in s with $(g - k)$ gaps
- Size of gapped match set is $O(g^{g-k})$, independent of $|\Sigma|$



(g,k)-Gappy Kernel

- Feature map:

For a g-mer s , $F_{(g,k)}(s) = (F_t(s))_{\{k\text{-mers } t\}}$,

where $F_t(s) = 1$ if t is in set $G_{(g,k)}(s)$,

$F_t(s) = 0$ otherwise;

extend additively by summing over g-mers s in x

AKQKL \longrightarrow (0 , ... , 1 , 1 , ... , 1 , ... , 1 , ... , 0)
AKK AKL AKQ AQK

- Weighted version with gap penalty, $0 < \gamma \leq 1$:

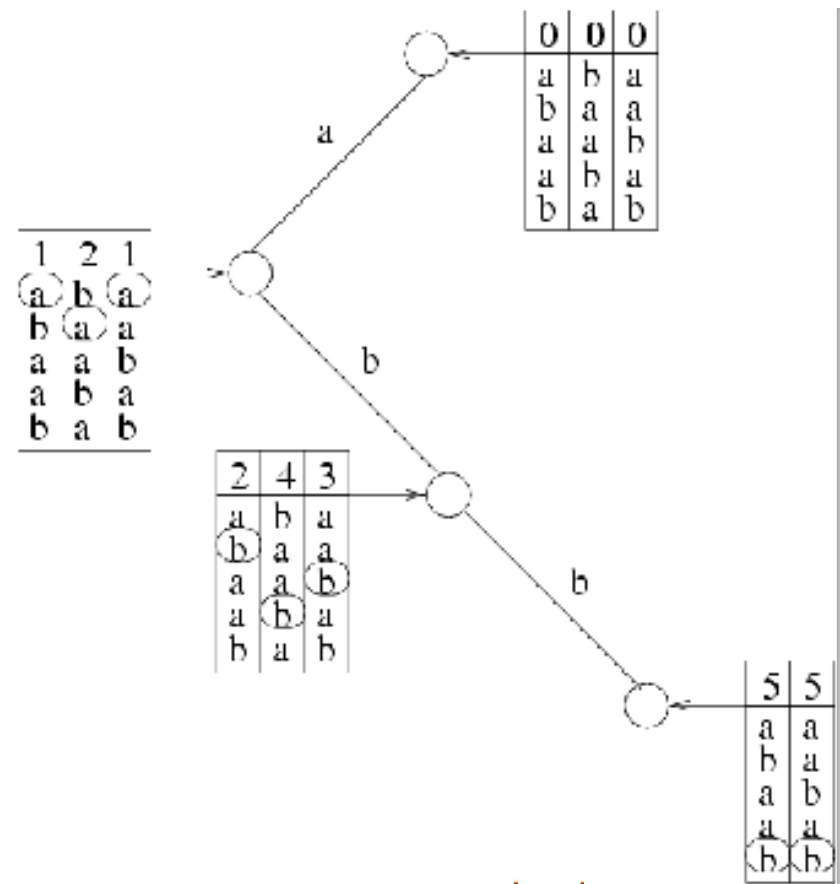
$F_t(s) = (1/\gamma^k) \sum_{\{\text{subseq}(s) = t\}} \gamma^{\text{length}(\text{subseq}(s))}$,

can be computed by dynamic programming

Gives truncated version of Lodhi *et al.* string kernel


Gappy Kernel Computation

- Traverse instance g-mers in the data, greedily align to k-length paths (k-mer features)
- At leaf node, count instances for each input sequence (or perform restricted dynamic programming for weighted version)



$$\square O(c_K(|x| + |y|)) \text{ with } c_K = g^{g-k+1}$$

Gappy Kernel SCOP Results



Inexact Matching through Probabilistic Substitutions

- Use *substitution matrices* to obtain $P(a|b)$, substitution probabilities for residues a, b
- The *mutation neighborhood* $M_{(k, \epsilon)}(s)$ is the set of all k -mers t such that

$$-\sum_{i=1 \dots k} \log P(s_i|t_i) < \epsilon$$


For a k -mer s , map $F_{(k, \epsilon)}(s) = (F_t(s))_{\{k\text{-mers } t\}}$,
where $F_t(s) = 1$ if t is in neighborhood $M_{(k, \epsilon)}(s)$,

$F_t(s) = 0$ otherwise;

extend additively

□ $c_K = k N_\epsilon$, where N_ϵ is max size
of mutation neighborhood

Substitution Kernel SCOP Results



Inexact Matching through Wildcards

- Introduce wildcard character “ \square ”, define feature space indexed by k-mers from $\Sigma^k \setminus \{\square\}$, allowing up to m wildcards

For a k-mer s , $F_{(k,m)}(s) = (F_t(s))_{\{k\text{-mers } t\}}$,
 where $F_t(s) = \square^{\text{\#wildcards in } t}$, if t matches s ,
 $F_t(s) = 0$ otherwise;

extend additively


AKQ \longrightarrow (0 , ... , 1 , ... , \square , ... , \square , ... , \square , ... , 0)

AKQ
AK \square
A \square Q
 \square KQ

\square Using pruned depth k trie over $\Sigma^k \setminus \{\square\}$,

$$c_K = k^{m+1}$$

Wildcard Kernel SCOP Results



Related Recent String Kernel Work

- For exact matching case, Vishwanathan and Smola compute *convex combinations* of kernels using suffix trees
- Ben-Hur *et al.* define a motif kernel: features are known motifs, stored using trie
- Li and Noble use feature vectors of pairwise alignment scores (Smith-Waterman, BLAST)
- Can describe all the kernels here using *transducer formalism* (finite state automata) of Cortes *et al.*

Outline

1. Protein classification: Mismatch kernel
 - SVMs and kernel methods
 - Inexact matching through mismatches
 - Efficient kernel computation, fast prediction
2. Experimental results on SCOP dataset
3. Other models for inexact matching
 - Kernels from gaps, substitutions, wildcards
4. Cluster Kernels: Semi-supervised methods
 - Using unlabeled data to change the kernel

Use of Unlabeled Data

- About 30,000 proteins with known structure (labeled proteins), but about 1 million sequenced proteins
- BLAST, PSI-BLAST: widely-used heuristic alignment-based sequence similarity scores
 - Good *local similarity score*, less useful for more remote homology
 - BLAST/PSI-BLAST E-values give good measure of distance between sequences
- Can we use unlabeled data, combined with good local distance measure, for *semi-supervised* approach to protein classification?

Cluster Kernels


- Use unlabeled data to change the (string) kernel representation
- *Cluster assumption*: decision boundary should pass through low density region of input space; clusters in the data are likely to have consistent labels
 - Profile kernel
 - Bagged kernel

J. Weston, C. Leslie, D. Zhou, A. Elisseeff, and W. S. Noble,
Cluster Kernels for Semi-supervised Protein Classification.
Submitted.

Profile Kernel

- Represent sequence x by the average sequences in its neighborhood $N(x)$:

$$F^{\text{Profile}}(x) = (1/|N(x)|) \sum_{x' \in N(x)} F(x')$$

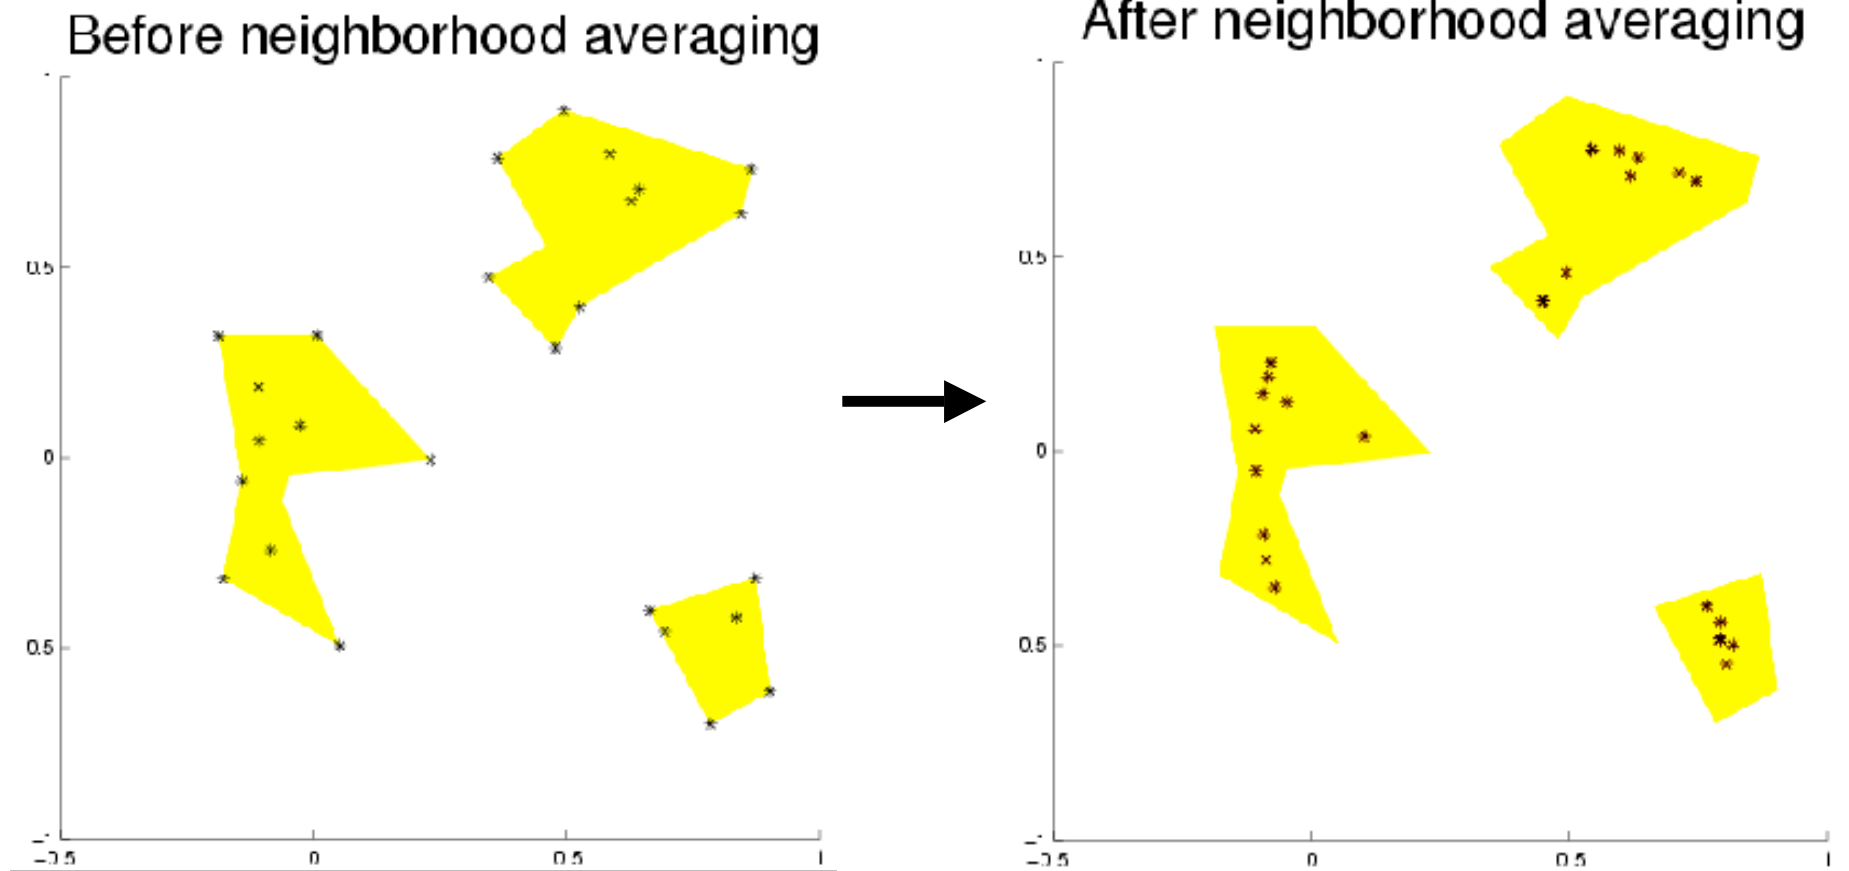


- Profile kernel:

$$K^{\text{Profile}}(x, y) = (1/|N(x)||N(y)|) \sum_{N(x) \times N(y)} K(x', y')$$

- Use PSI-BLAST distance and mismatch kernel as base kernel

Profile kernel addresses cluster assumption



Bagged Kernel

- Use k-means clustering to cluster data (labeled + unlabeled), N bagged runs

- Using N runs, define

$$p(x,y) = (\# \text{ times } x, y \text{ are in same cluster})/N$$

- Bagged kernel:

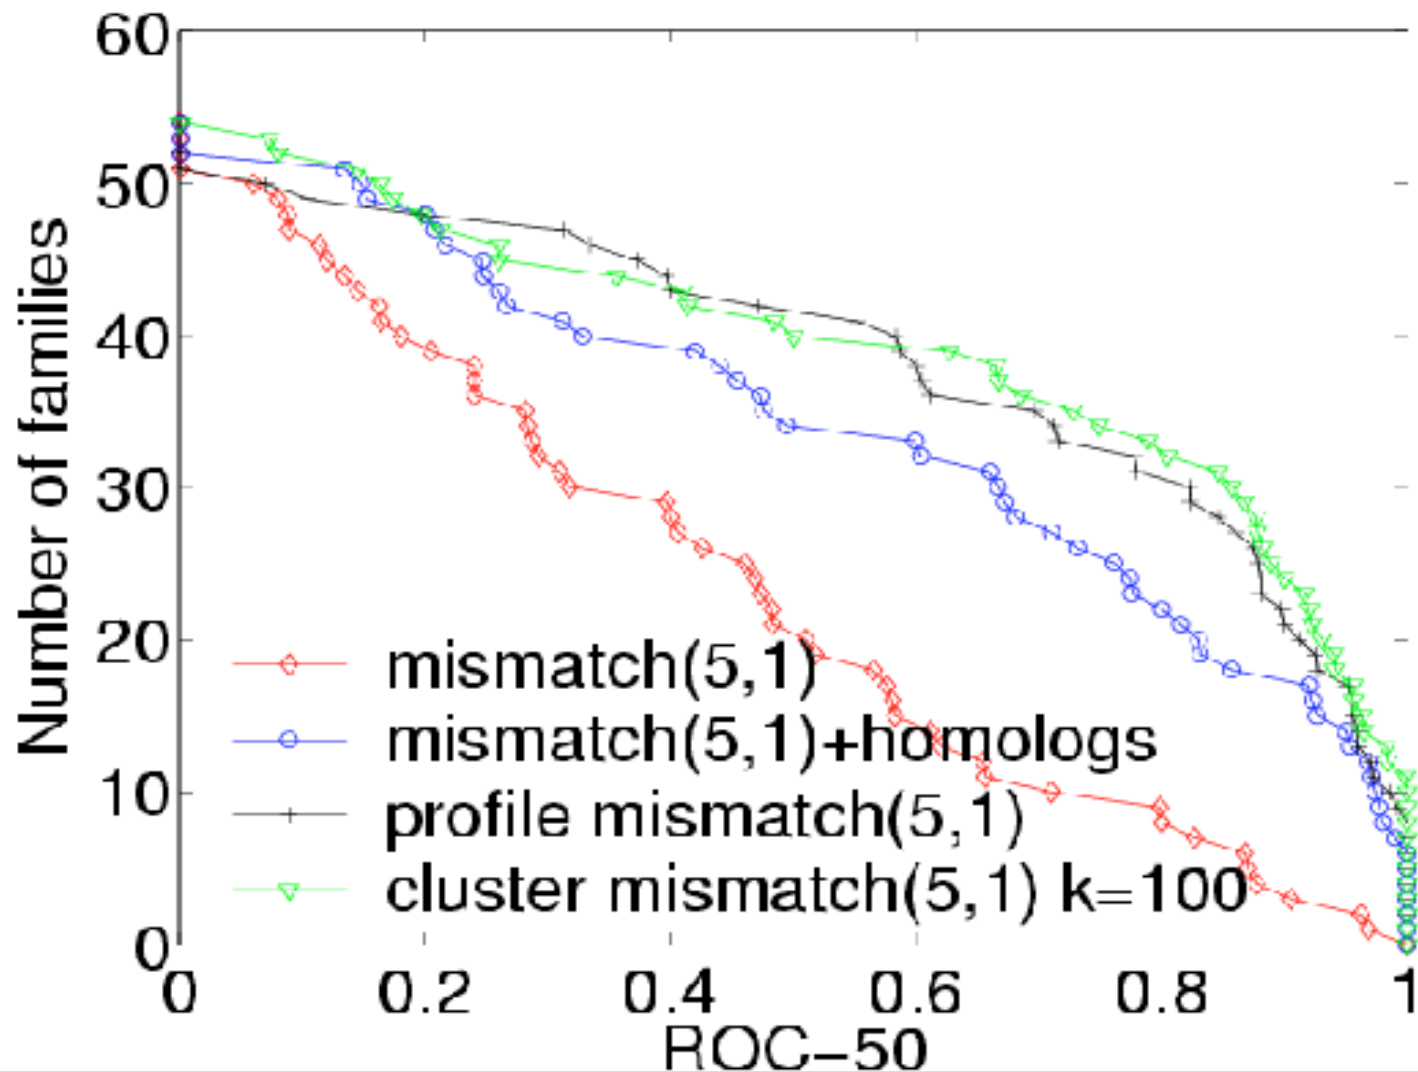
$$K^{\text{Bagged}}(x,y) = p(x,y) K(x,y)$$

- Use PSI-BLAST for clustering, mismatch kernel for underlying kernel

Experimental Set-up

- *Full dataset:* 7329 SCOP protein sequences
- *Experiments:*
 - 54 target families (remote homology detection)
 - Test + training approximately for each experiment <4000 sequences, other data treated as unlabeled
- *Evaluation:* How well do cluster kernel approaches compare to the standard approach, adding positive homologs to dataset?

Results for Cluster Kernels



Conclusions

- SVMs with *string kernels* – like mismatch kernels – that incorporate *inexact matching* are competitive with best-known methods for protein classification
- Efficient kernel computation: $O(c_K(|x| + |y|))$, linear-time prediction, feature extraction
- Gaps, substitutions, and wildcards give kernel constant c_K that is independent of alphabet size
- Semi-supervised *cluster kernels* – using unlabeled data to modify kernel representation – improve on original string kernel

Future Work

- Full *multiclass* protein classification problem
 - 1000s of classes of different sizes, hierarchical labels
 - Use of unlabeled data for improving kernel representation
- *Domain segmentation* problem: predict and classify domains of multidomain proteins
- Develop and implement *semi-supervised ranking* approaches, make available on web server
- *Local structure* prediction: predict local conformation state (backbone angles) for short peptide segments, step towards structure prediction


Protein Ranking

- *Pairwise sequence comparison*: most fundamental bioinformatics application
- BLAST, PSI-BLAST: widely-used heuristic alignment-based sequence similarity scores
 - Given query sequence, search unlabeled database and return ranked list of similar sequences
 - Query sequence does not have to belong to known family or superfamily
 - Good local similarity score, less useful for more remote homology
- Can we use *semi-supervised machine learning* to improve on PSI-BLAST?

Joint work with J. Weston, A. Elisseeff, and W. S. Noble

Ranking Induced by Clustering

- *Idea:* Map out protein sequence space by performing constrained clustering, using label constraints




- Output ranking: rank by cluster, nearest cluster first
- Use dissimilarity measure derived from PSI-BLAST
- Best to use constrained clustering for model selection (number of clusters) based on labeled data, then use regular efficient clustering algorithms
 - Generalized (“kernel”) k-means
 - Hierarchical clustering (average linkage)

Experimental Set-up for Ranking


- *Training set:* 4246 SCOP protein sequences (from 554 superfamilies) – known classes
- *Test set:*
 - 3083 SCOP protein sequences (from 516 different superfamilies) – hidden classes
 - 101,403 unlabeled sequences from SWISSPROT
- *Task:* How well can we retrieve database sequences (from train + test sets) in same superfamily as query? Evaluate with ROC-50 scores
- For initial experiments, SWISSPROT sequences only used for PSI-BLAST scores, not for clustering

Ranking Results for Clustering



Label Propagation


- Zhu and Ghahramani: Propagate labels through dense regions of example space
 - Y is $m \times 2$ matrix of label probabilities, where m is number of examples
 - Clamp known labels:
(1,0) for class 1, (0,1) for class 2
 - K is matrix of transition probabilities (sparse, derived from PSI-BLAST)
- Iterate until convergence to fixed point:
- Propagate: $Y \leftarrow K Y$
 - Row normalize Y
 - Clamp known labels




Online Semi-Supervised Approach

- *Idea*: PSI-BLAST returns good high-confidence prediction scores, can rule out extremely low-confidence scores
 - Given query, assign *positive pseudo-label* to sequences with good (small) E-values
 - Assign *negative pseudo-label* to sequences with poor (large) E-values
- Small number of (pseudo-)labeled examples, rest of database considered unlabeled: apply semi-supervised technique
- Known class information not used, but can use for model selection


Ranking Results for Label-Prop




Ranking Results: Comparison with Structure-Structure Scores




(5,1)-Mismatch vs Fisher Using ROC Scores



(5,1)-Mismatch vs Fisher Using ROC-50 Scores



(5,1)-Mismatch vs. 3-Spectrum Using ROC Scores



(5,1)-Mismatch vs. 3-Spectrum Using ROC-50 Scores

